

JARED MONDSCHIEIN

Contingency Frameworks for Future U.S.-China Cooperation on AI Assurance and Security

Preserving Strategic Optionality

The U.S.-China relationship is the most important international relationship for the future of artificial intelligence (AI). Combined, the two countries' AI ecosystems incorporate the vast majority of AI researchers, deploy nearly all AI compute capacity, and absorb the bulk of global AI investment.¹ American and Chinese companies have released the most powerful and popular foundation models and AI systems and are creating global dependencies on American and Chinese AI technologies.²

These global dependencies create significant leverage for the United States and China. American and Chinese standards and norms for technical assurance, reliability, and security will become the global baseline. Operationalizing this leverage to secure American advantage has evolved into a priority objective of the second administration of President Donald Trump. Indeed, President Trump's AI Action



Plan seeks to establish American AI as the standard for global AI capabilities and aims to ensure that the global AI build-out is constructed on the U.S. AI technology stack, from chips to applications, and is aligned with U.S. values and interests.³ To do this, the Trump administration could, and has already begun to, take a series of unilateral actions and lead bilateral and multilateral initiatives with U.S. allies and partners.

Opportunities for U.S.-China AI cooperation could emerge within this competitive context. Collaboration is already ongoing, including academic partnerships,⁴ commercial cooperation,⁵ and Track 2 dialogues.⁶ Future bilateral engagements on AI assurance, reliability, and security may be beneficial to U.S. technology leadership, economic security, and national security.⁷

This paper explores how the United States and China might structure such cooperative efforts should such cooperation become advantageous. I describe five lines of effort (LOEs) grounded in national interests that could foster a stable, predictable international environment for AI

development and deployment. While current geopolitical competition, deep mutual distrust, and the strategic stakes involved may make bilateral AI cooperation neither feasible nor sensible in the near term, this analysis examines what frameworks and mechanisms would be required if circumstances changed. Rather than advocating for immediate cooperation, this paper aims to inform contingency planning under the premise that preserving strategic optionality for future engagement serves both nations' interests. Understanding how cooperation could work—even if it proves unattainable—can help policymakers preserve flexibility in an uncertain technological landscape.

Extant U.S.-China Agreements

As of February 2026, the United States and China have both been parties to several nonbinding multilateral agreements on AI.⁸ These include

- **the Bletchley Declaration**, which crafted an international agenda for fostering public trust and confidence in AI systems by addressing frontier AI risks⁹
- **United Nations General Assembly Draft Resolution A/78/L.49, “Seizing the Opportunities of Safe, Secure and Trustworthy Artificial Intelligence Systems for Sustainable Development,”** a U.S.-led resolution primarily focused on (1) generating international cooperation to leverage safe, secure, and trustworthy AI to advance the United Nations’ Sustainable Development Goals and (2) managing the risks of AI to sustainable development and human rights¹⁰

Abbreviations

AI	artificial intelligence
AISEA	AI Safety and Efficacy Administration
ASI	artificial superintelligence
FAIA	Federal AI Administration
FLOPS	floating-point operations per second
IYAI	International Year of Artificial Intelligence
LOE	line of effort
OECD	Organisation for Economic Co-operation and Development
R&D	research and development

Rather than advocating for immediate cooperation, this paper aims to inform contingency planning under the premise that preserving strategic optionality for future engagement serves both nations' interests.

- **United Nations General Assembly Resolution 78/311, “Enhancing International Cooperation on Capacity-Building of Artificial Intelligence,”** a China-led resolution to boost international financial and technical assistance and capacity-building efforts to strengthen the developing world’s ability to design, develop, deploy, decommission, and use AI capabilities.¹¹

These three instruments are nonbinding political statements: They endorse high-level principles and encourage risk management and capacity-building but create no enforceable limits on areas of potential risk, such as accelerated AI research and development (R&D) or military use. They call for safeguards and dialogue, avoid bans on specific military applications, and leave implementation to national policy and voluntary cooperation.

In aggregate, these instruments suggest that the United States and China accept that AI presents both economic opportunities that could be accelerated and significant risks that could be managed via international cooperation. Both countries were willing to agree that AI development and use should respect human dignity, pro-

mote sustainable development, and avoid undermining international peace and security. While the Trump administration’s vision for international AI policy is still taking shape, the AI Action Plan indicates broad continued U.S. support for these pillars.¹² Chinese government statements likewise indicate support for continued diplomatic engagement in these domains.¹³

Proposed AI Treaties Go Further Than Current Agreements

Authors of proposed international AI agreements analogize their proposals to treaties governing nuclear weapons. However, AI technologies exhibit fundamental differences from nuclear weapons that require a requisite shift in assumptions and framing (Table 1).

I reviewed three proposed international AI agreements to demonstrate the pitfalls of this analytic trap:¹⁴ the Treaty on the Prohibition of Dangerous Artificial Intelligence,¹⁵ the Treaty on Artificial Intelligence Safety and Cooperation,¹⁶ and the International Agreement to Prevent the Premature Creation of Artificial Superintelli-

TABLE 1
Comparison of Characteristics of Nuclear Weapons and AI

Characteristic	Nuclear Weapons	AI
Consensus of catastrophic risk	There is an international consensus that use would be catastrophic, mutual vulnerability is real, and some domains create acute escalation and proliferation risks.	There is no clear model of risk. Foci include existential risk from frontier models; near-term economic harms; and missed economic, social, and military opportunities.
Political framing	Catastrophes created coalitions of stakeholders who generated political will for mutual restraint and scientific cooperation.	Private-sector stakeholders and development coalitions emphasize economic and development opportunities.
Chokepoints	Nuclear weapons are centrally located, owned by governments, and directly correlated with demonstrable risks (e.g., fissile material, delivery systems).	The AI technology stack is relocatable, owned by private-sector entities, and loosely or indirectly correlated with potential risks (e.g., chips, manufacturing equipment, critical minerals, electricity).
Domains	Nuclear weapons have discrete, spatially bounded domains (e.g., test environments, Antarctica, orbit) and dual-use functions (e.g., energy).	General-purpose AI technologies are deeply embedded across the economic and military domains; risks could emerge via integration with military or civilian systems.
Capability symmetries	There are relatively symmetric superpower arsenals and clear mutual vulnerability.	There is an uneven distribution of capabilities, data, and industrial capacity across many countries.
Incentives	Unconstrained racing and proliferation increased self-risk of annihilation without proportional gain.	There is an expectation of ongoing economic and military gains from out-innovating rivals.

NOTE: This table summarizes the author's review of nuclear weapons treaties, including the Antarctic Treaty; the Treaty on Principles Governing the Activities of States in the Exploration and Use of Outer Space, Including the Moon and Other Celestial Bodies; the Treaty Banning Nuclear Weapon Tests in the Atmosphere, in Outer Space and Under Water; the Treaty on the Non-Proliferation of Nuclear Weapons; and the Strategic Arms Reduction Treaty. The assessments of AI characteristics were adapted from Horowitz and Kahn, "Nuclear Non-Proliferation Is the Wrong Framework for AI Governance."

gence¹⁷ (summaries are included in Table A.1 in the appendix). Table 2 highlights the results of my assessment of AI technologies (listed in Table 1) against aspects of these three proposed agreements.

These treaties' *raison d'être*, to mitigate the catastrophic risks that certain AI technologies pose to

humanity, constrains their ability to generate the pre-requisite political economies. The history of international agreements to control items of economic and/or military means strongly suggests that there must be a demonstrable realization of risk to generate sufficient political will to align national interests. Indeed, national regulations and

TABLE 2
Assumptions of Proposed AI Agreements Against Current AI Fundamentals

Characteristics	Assumptions of Proposed International AI Agreements	Aligned with Current Fundamentals of AI Technologies?
Consensus of catastrophic risk	Assumes widespread agreement that certain technological domains pose catastrophic risks to humanity	No
Political framing	Assumes coalition agreement that risk mitigation should outweigh national economic, military, and development ambitions	No
Chokepoints	Assumes government ability to control and monitor and/or have knowledge of AI training runs; assumes government ownership, monitoring, and/or control of AI chips and semiconductor manufacturing equipment; assumes that control over chips and equipment mitigates risk	No
Domains	Assumes that risk is derived from both civilian and military applications of AI	Yes
Capability symmetries	Assumes that AI capabilities and industrial capacity are evenly distributed among states	No
Incentives	Assumes that there is a political expectation that increased self-risk of catastrophe is not overcome by or correlated with expectation of outcompeting economic and military rivals	No

NOTE: This table summarizes the author's assessment of AI technologies listed in Table 1 against the Treaty on the Prohibition of Dangerous Artificial Intelligence (Miotti, "An International Treaty to Implement a Global Compute Cap for Advanced Artificial Intelligence"), the Treaty on Artificial Intelligence Safety and Cooperation (TAISC.org, "Treaty on Artificial Intelligence Safety and Cooperation"), and the International Agreement to Prevent the Premature Creation of Artificial Superintelligence (Scher et al., "An International Agreement to Prevent the Premature Creation of Artificial Superintelligence"). These treaties are further summarized in Table A.1 in the appendix.

their related international agreements for controlling civilian, military, and dual-use technologies (e.g., pharmaceuticals, food, transportation, telecommunications, manufacturing, nuclear weapons, biological and chemical weapons) are typically written in the aftermath of catastrophes. As of this writing (spring 2026), there is no widespread consensus among technical experts and policymakers of the types, severity, and pathways of AI risks.¹⁸

Cooperation to restrict nuclear weapons grew from foundations of scientific cooperation. Historic examples of (somewhat)¹⁹ proactive agreements between adversaries, like the Antarctic Treaty and the Treaty on Principles Governing the Activities of States in the Exploration and Use of Outer Space, Including the Moon and Other Celestial Bodies,²⁰ focus on furthering scientifically promising, politically acceptable means of scientific cooperation while

reducing potential military and political impediments. The proposed AI agreements are a stark departure from other preemptive international agreements, including recent multilateral AI agreements to which China and the United States are parties that seek out areas of constructive cooperation to unlock, not block, economic and technological AI goals.

Preserving Optionality for U.S.-China AI Engagement

Overly restrictive proposals are often misaligned with the fundamentals of AI technologies and/or are geopolitically naïve. They harm the effort to generate meaningful international cooperation between strategic rivals. They also risk alienating both companies and governments, fostering and solidifying resistance to future opportunities.

Instead, a realistic approach should expect the United States and China to act in their national interests and operate within the boundary conditions imposed by their worldviews, national AI strategies, political economics, and international geopolitics:

- **Assumption 1:** The governments of the United States, China, and other countries view AI as a key enabler for achieving broad economic and military benefits.
- **Assumption 2:** Governments and AI companies believe that the “winners” of the AI race will reap a disproportionate share of the benefits of AI technologies.
- **Assumption 3:** The private sector will retain significant ownership and control over AI technologies.
- **Assumption 4:** Political leadership will be reacting to emerging economic and national security risks of AI as they emerge and will focus proactive policies toward maximizing economic and national security benefits.
- **Assumption 5:** U.S.-China competitive and cooperative actions will take place within a context in which (1) the predominant view of Chinese political leadership is that the United States and its allies and partners are determined to contain, suppress, and encircle China and (2) the predominant view of U.S. political leadership is that China is determined to displace the United States as the world’s economic and military leader.
- **Assumption 6:** U.S. and Chinese leadership are constrained by domestic political economies that discourage legally binding bilateral treaties. Any path forward must begin with nonbinding principles and technical collaboration.
- **Assumption 7:** U.S., Chinese, and other firms are deeply entangled in AI supply chains, which are global in scope, and complete decoupling is unrealistic.
- **Assumption 8:** The Chinese Communist Party’s worldview is dominated by a pragmatic realist focus on internal continuity. The current worldview of American political leadership aims to maximize financial gains while minimizing government interventions outside of sector-specific controls.

My consideration of these boundary conditions yields unilateral, bilateral, and multilateral confidence-building measures, or LOEs, that could create a stable, predictable international environment in which trustworthy AI

systems can be developed, deployed, and used. Each LOE reinforces the next, transforming abstract cooperation into a self-reinforcing architecture of trust.

LOE 1, building international consensus on AI risks, creates the evidentiary foundation and political awareness necessary for LOE 2's reframing of AI from competition to cooperation. Those two steps, in turn, generate the legitimacy needed for LOE 3's formation of chokepoints and oversight institutions, which require both technical standards and political buy-in to function. Narrow domain-specific agreements in LOE 4 depend on demonstrable success in LOE 3 and early cooperative habits developed through test ranges or joint investigations; these agreements serve as pilots, feeding lessons into broader symmetry-building efforts under LOE 5. Only after these technical, institutional, and relational preconditions are met can aligned political, economic, and security incentives produce a stable environment for sustained interaction.

Lines of Effort

LOE 1: Generate International Consensus on AI Risks

Similar to what the Franck Report authors articulated in 1945 about nuclear arms,²¹ my position is that only tangible demonstrations of AI risks will mobilize global leadership toward risk-reduction measures. As the Franck Report suggested, this can be achieved via controlled demonstrations or through real-world employment and catastrophe that amplify Track 2 dialogues and private-sector exchanges. Laboratory-based experiments must be genuinely scientific, with opportunities for technical expert engagements; polit-

Only tangible demonstrations of AI risks will mobilize global leadership toward risk-reduction measures.

ical usurpation would threaten their credibility. Experiments must also be shaped to mitigate concerns around the sharing of proprietary information. Reactive cooperation following real-world incidents must reckon with political economies that often favor amplified competition and mistrust in the aftermath.

LOE 2: Shift the Political Framing

Current American and Chinese political leadership view AI almost exclusively through a lens of economic development and military opportunities. Diverse coalitions of stakeholders should work to shift toward a rationale grounded in the benefits of international cooperation, particularly in the scientific and development domains. Science and technology cooperation has long been recognized as an avenue to strengthen geopolitically challenging partnerships.²² International efforts should work toward creating coalitions of national governments, AI firms, critical infrastructure providers, technical and scientific institutions, organized labor, civil society, and multilateral

bodies and standard-setting organizations.²³ This could mirror the International Geophysical Year, a 1957–1958 international effort of intensive geophysical observations and measurements.²⁴

LOE 3: Create and Identify Chokepoints

AI-enabling inputs—chips, semiconductor manufacturing equipment, etc.—are often identified as focal points for national and international control. However, the constraining capacity of AI supply chain inputs is questionable because of their high mobility, rapidly increasing algorithmic efficiencies, and the global diffusion of these inputs. Indeed, the proposed AI agreements summarized above fall victim to this analytic trap.

Other technologies—pharmaceuticals, aviation, etc.—have encountered similar obstacles and created a common solution: employing market access as the chokepoint. I propose to adopt this model and employ U.S. and global market access as a chokepoint. This model would leverage U.S. market power to enforce risk-reduction standards and assert U.S. prerogatives that align international standards to U.S. requirements. Multilateralizing this approach via the Organisation for Economic Co-operation and Development (OECD) or International Organization for Standardization safety standards adoption could amplify benefits and reduce conditions for potential retaliations.

LOE 4: Achieve Narrow Agreements in Tightly Scoped Domains

Focusing on tightly scoped agreements within discrete AI sectors that can be carved out from economic and military ambitions creates political economies for agreements. These could serve as confidence-building pilots leading to broader agreements on norms, safeguards, and responsibilities. Symbolic agreements could accelerate shifts in rhetoric even without enforcement. See, for example, the Biden-Xi joint statement affirming the need to maintain human control over the decision to use nuclear weapons.²⁵

LOE 5: Create and Preserve Symmetries

The global AI ecosystem is currently characterized by an imbalanced distribution of capabilities, data, and industrial capacity. This imbalance, combined with a general belief that AI implementation will deliver uneven benefits to a select few, creates a context in which any restrictions will generate opposing reactions by those perceived to be disadvantaged. Two pathways emerge: (1) AI leaders dedicate real technical, financial, and diplomatic resources to boost national AI capacities (e.g., compute capacity, data access, human capital) to a global baseline that allows nearly every country to capitalize on AI opportunities, believe it is being treated fairly, and believe it will not be left behind in an AI world, or (2) AI leaders embrace the asymmetry of AI and maximize differential capacity so as to confine the economic and military benefits (and risk management) of AI to a select few countries.²⁶

Conclusion

AI is one of many challenges that have become irritants in the U.S.-China relationship. I expect both countries to act in their national interests and operate within the boundary conditions imposed by worldviews, national AI strategies, technoeconomics, and international geopolitics. While today's political economies may not support bilateral engagement, contingency planning builds in optionality to enable flexible policy responses as circumstances change. The proposed roadmap of LOEs seeks to employ confidence-building measures to align political, economic, and security incentives. I believe that the

combination of an international consensus on AI risks, a shift in the political framing around AI toward scientific and economic cooperation, the identification of useful chokepoints, agreements within tightly scoped economic or geographic domains, and the deliberate creation of symmetries (or asymmetries) can align the political, economic, and security incentives of AI. As the United States and China race toward transformative AI capabilities, maintaining the option for future coordination—however remote that possibility may seem today—represents prudent geopolitical strategy for navigating an unprecedented technological transition.

While today's political economies may not support bilateral engagement, contingency planning builds in optionality to enable flexible policy responses as circumstances change. The proposed roadmap of LOEs seeks to employ confidence-building measures to align political, economic, and security incentives.

APPENDIX

Example International Engagement Activities and Summary of Proposed AI Agreements

This appendix contains broad concepts for U.S.-led international efforts on AI assurance, reliability, and security that could create a stable, predictable international environment for AI development, deployment, and use.²⁷ The intent is to provide initial suggestions for contingency planning purposes should such cooperation become advantageous for U.S. technology leadership, economic security, and national security. Within each LOE, some activities may conflict with each other and others could be sequenced to reinforce each other, reflecting the intent to provide optionality that can flex with the political environment. These activities could be implemented as a mixture of multilateral and bilateral efforts with varying degrees of Chinese government participation, as appropriate.

Activities for LOE 1: Generate International Consensus on AI Risks

International AI test ranges: International AI test ranges would be jointly governed, secure environments in which countries, companies, and researchers could conduct controlled demonstrations and rigorous evaluations of AI systems. These ranges would combine high-assurance technical infrastructure, red-teaming capabilities, and scenario design (for example, in cyber-, bio-, financial-, and military-relevant domains) to safely probe how powerful models behave under stress, including in edge cases and adversarial conditions. By generating shared, indepen-

dently validated evidence on both capabilities and failure modes, international AI test ranges would help build a common fact base on AI risks; reduce mistrust among major powers; and support the development of concrete international norms, safeguards, and crisis-management mechanisms before dangerous capabilities are widely deployed. Participation could be structured through an intergovernmental organization or public-private partnership in which member states co-fund, co-govern, and share access to the facilities and their scientific outputs on a long-term basis.

Precedents: International Thermonuclear Experimental Reactor, European Organization for Nuclear Research, North Atlantic Treaty Organization Cyber Range.

Reaction to an incident: A catastrophic AI incident would create a moment in time when abstract concerns become concrete, making it easier to align governments around shared interests in risk reduction.²⁸ In the immediate aftermath, the United States could push for a joint, multinational technical investigation that produces an authoritative, public record of what went wrong; which specific AI capabilities and governance failures were implicated; and how similar failures could recur across jurisdictions. Coupling that common fact base with high-level political forums (e.g., the United Nations, the Group of 20, a dedicated summit) in which leaders endorse the technical findings and acknowledge that no state can unilaterally manage such risks can reframe the incident from a one-off tragedy into a systemic warning signal. That, in turn, creates political space to codify minimum global safeguards—such as mandatory incident reporting, shared red-team and evaluation protocols, and baseline controls on the most-

dangerous training runs or applications—because they can be explicitly justified as direct lessons learned from a real disaster rather than hypothetical scenarios.

Precedents: Air accident inquiries, joint responses to terrorist attacks.

Activities for LOE 2: Shift the Political Framing

International Year of Artificial Intelligence (IYAI): IYAI would be a time-limited, globally coordinated scientific campaign in which countries, companies, and researchers jointly study and stress-test advanced AI systems using shared methods and open results, creating a common empirical foundation for governance. Over 18–24 months, participants would run synchronized evaluations of AI systems’ safety, robustness, and misuse potential; build global observation networks for AI impacts and incidents; and carry out coordinated experiments on AI’s economic and societal effects, with anonymized data and findings made widely available. The effort would also pilot practical norms on transparency, data-sharing, and safety disclosures and result in permanent infrastructure—such as an international AI observatory, open benchmark repositories, and expert groups on high-risk domains—to inform future rules, standards, and agreements on AI systems.

Precedents: International Geophysical Year, 2025 International Year of Quantum Science and Technology.

Global AI alliance: This coalition would maximize the upside of AI technologies by pooling capabilities, markets, and expertise under shared rules that make powerful systems broadly beneficial and widely trusted. Instead of

each country and company racing alone, members would coordinate investments in research, compute infrastructure, and open scientific resources so that breakthroughs in such areas as health, climate, and productivity diffuse quickly across the alliance. Governments would align their regulations, standards, and incentives to make it easier for responsible firms to scale AI solutions across jurisdictions, while joint safety and evaluation frameworks reduce duplication and unlock cross-border deployments that would otherwise be blocked by mistrust. By bringing in labor, civil society, and developing countries as real rule-shapers and by using finance, insurance, and liability tools to reward low-risk, high-impact innovation, the alliance would aim to turn responsible AI into a competitive advantage—growing economies, strengthening democracies, and expanding access to beneficial AI systems, even as it quietly keeps the most hazardous capabilities under common safeguards.

Precedents: OECD Global Partnership on AI, Pax Silica.

Activities for LOE 3: Create and Identify Chokepoints

AI Safety and Efficacy Administration (AISEA): The AISEA would be a U.S. federal entity with authority for oversight of AI technologies modeled on the rigor of pharmaceutical regulation. The AISEA would evaluate AI systems before they were deployed or exported, requiring safety testing, transparency documentation, and periodic inspections to ensure compliance with federal standards. AISEA approvals of foreign AI technologies would be required prior to export to the United States. Like the U.S. Food and Drug Administration (FDA), this lever would use

U.S. market access to encourage other countries to develop similar standards and institutions.

Precedent: FDA.

Federal AI Administration (FAIA): The FAIA would be modeled on the Federal Aviation Administration (FAA) and would focus on certifying and supervising real-world AI operations. The FAIA would license AI “operators” (such as major developers and deployers); certify AI systems and classes of systems for specific use cases; and publish detailed “rules of the air” governing how AI technologies can be built, integrated, and deployed in safety-critical and high-impact environments. It would require continuous monitoring; mandatory incident and near-miss reporting; and the authority to ground or restrict systems when safety, reliability, or security concerns arise. It would also update technical and operational standards as the broader AI ecosystem evolves. The FAIA would work with international counterparts to create globalized standards, approve export destinations for U.S. AI technologies, and certify countries as accredited exporters of AI technologies to the United States.

Precedent: FAA.

Activities for LOE 4: Achieve Narrow Agreements in Tightly Scoped Domains

Safeguarding the nuclear threshold: Nuclear powers would pledge not to delegate nuclear launch authority to AI and not to use AI systems—directly or via cyber operations—to disrupt or spoof each other’s nuclear command-and-control or early-warning networks. The agreement would be incorporated as an addendum to the Treaty on the

Non-Proliferation of Nuclear Weapons and implemented through a standing technical working group on human control in nuclear decisionmaking and dedicated crisis-communication procedures for AI-related anomalies in the space and cyber domains. Initially framed as a political (not legally binding) commitment with unilateral transparency measures, it could later evolve into verifiable, tightly defined bans on specific AI uses against nuclear command-, control-, and communications-related assets.

Precedents: U.S.-China statement pledging to maintain human control over nuclear launch authority, Strategic Arms Limitation Talks I limitations on interference with national technical means of verification.

Blocking AI-enabled pathogen design: Governments would agree to prohibit training, fine-tuning, or releasing models that provide actionable, step-by-step assistance for creating novel pathogens, optimizing their spread, or bypassing detection and response; they would also bar each side from deliberately exploiting such models for offensive bio purposes. Parties could also collaborate on shared safeguards—such as red-teaming protocols, access controls, and content filters for high-risk bio capabilities—paired with an incident-notification mechanism if dangerous bio-relevant AI behaviors are discovered. This would become a negotiated addendum to the Biological Weapons Convention.

Precedents: Biological Weapons Convention; 1975 Asilomar Conference on Recombinant DNA.

U.S.-China Basic Principles of AI Relations agreement: A U.S.-China Basic Principles of AI Relations agreement would be a short political agreement that sets shared

guardrails for how both sides develop and use advanced AI in ways that affect strategic stability, crisis management, and global security. It would affirm that neither country will use AI systems to interfere with the other's nuclear command and control; neither will delegate nuclear launch authority to AI; and neither will conduct dangerously escalatory autonomous AI operations in the air, maritime, space, or cyber domains, while recognizing a mutual interest in mitigating and preventing accidents, miscalculation, and unauthorized AI behavior. The agreement would endorse principles already agreed to in multilateral agreements, including those at the intersection of AI use and human rights, implementation of regulatory and governance approaches, and strengthening capacity-building activities with developing countries.²⁹

Precedent: 1972 U.S.-USSR Basic Principles agreement.

Activities for LOE 5: Create and Preserve Symmetries

AI for all: The governments of the United States, China, and other leading AI nations; companies; universities; civil society; and multilateral institutions would agree to launch a sustained effort to ensure that every country can meaningfully participate in an AI-driven global economy. Investments in technical infrastructure, talent development, and shared research platforms coupled with direct financial support would aim to help all countries achieve a baseline capacity for AI innovation, adoption, and use. This global standard would be sufficient to accelerate economic gains while limiting military advantage.

Precedents: Partnership for Global Inclusivity on AI, Atoms for Peace.

Unilateral AI hegemony: The United States could unilaterally decide to restrict all access to AI technologies. Exports of any AI-relevant technology (semiconductor manufacturing equipment, chips, models, etc.) would require a license; applications would be reviewed on a case-by-case basis for a select few countries and under a presumption of denial for all others. Access to U.S.-hosted cloud compute and foreign-hosted cloud compute employing U.S. technologies would be restricted to U.S. entities. Within the United States, AI research would be considered “born classified” and restricted to a select few institutions. The United States could take unilateral military or economic action to prevent unintended proliferation of AI capabilities.

Precedent: U.S. export controls.

Shared AI hegemony: The United States, China, and other leading AI nations would agree to allow access to AI technologies only to parties of the agreement. The parties would take steps to restrict AI exports and access to cloud compute and research. The parties would agree to take mutual action to prevent proliferation of AI capabilities beyond their membership. Finally, the parties would agree to mutually beneficial AI governance and unrestricted access to technologies and information.

Precedent: Wassenaar Arrangement.

Table A.1 summarizes three proposed AI agreements.

TABLE A.1
Summary of Proposed AI Agreements

	Treaty on the Prohibition of Dangerous Artificial Intelligence ^a	Treaty on Artificial Intelligence Safety and Cooperation ^b	International Agreement to Prevent the Premature Creation of Artificial Superintelligence ^c
Goal is to	Reduce risks from the development of advanced AI.	Prevent the development of AI systems that pose existential risks to humanity.	Prevent the premature development of artificial superintelligence (ASI) until AI development can proceed without posing catastrophic risks.
Assumes that	1. Premature development of ASI poses catastrophic risks that outweigh beneficial AI capabilities, American AI leadership, intellectual property leaks, and costs to the private sector. 2. Buying time provides space to solve for alignment to human values, geopolitical destabilization, mass AI-caused unemployment, and misuse of powerful AI. 3. Floating-point operations per second (FLOPS) are a reliable indicator of risk.		
Prohibits	1. Civil or military development, deployment, transfer, possession, and use of AI above a moratorium threshold (10^{24} FLOPS) 2. Development and use of human-level AI, artificial general intelligence, or ASI for civilian or military uses	Training of AI models greater than 10^{23} FLOPS	1. Development or deployment of ASI and access to AI chips or ASI-relevant manufacturing capabilities to non-parties 2. Training runs exceeding strict threshold (10^{24} FLOPS) or post-training runs exceeding strict post-training threshold (10^{23} FLOPS) 3. Production of AI chips at facilities where tracking and monitoring is not feasible or implemented 4. R&D for (1) improvements for training general-purpose AI systems that would significantly increase model capabilities or reduce computational resources required to develop such systems, (2) distributed or decentralized training methods, (3) advancements in AI chip fabrication, and (4) design of more-performant or -efficient AI chips

Table A.1—Continued

	Treaty on the Prohibition of Dangerous Artificial Intelligence ^a	Treaty on Artificial Intelligence Safety and Cooperation ^b	International Agreement to Prevent the Premature Creation of Artificial Superintelligence ^c
Parties must	1. Regulate the development and use above a danger threshold (10^{21} FLOPS). 2. Implement regulations to monitor AI development and reporting of the development or use of “widely dangerous” AI systems.	1. Dismantle or repurpose computing clusters capable of conducting a training run with a performance of 3×10^{17} FLOPS. 2. Submit monthly reports on AI developers, the amount of computation, and safety measures. 3. Make facilities and organizations available for compliance monitoring.	1. Assist or do not impede measures by other parties to dissuade or prevent development by and within non-party states and jurisdictions. 2. Report training runs above the monitored threshold (10^{22} FLOPS). 3. Seek approval for training runs above the monitored threshold. 4. Report and monitor covered chip clusters (clusters with >16 H100 equivalents). 5. Report any planned transfer of AI chips. 6. Make AI chip production facilities and chip production inputs available for monitoring. 7. Seek approval for sale or transfer of AI chips and AI chip manufacturing equipment (within or between parties has presumption of approval; AI chip manufacturing equipment within or between parties is case by case; sale or transfer to non-parties has presumption of denial). 8. Make covered chip clusters available for monitoring and verification of chip use. 9. Maintain awareness of and relationships with researchers and organizations.
Creates	1. An annual convening to review thresholds 2. An international agency to serve as Secretariat 3. An international government and private-sector hotline for AI security threats	1. An International AI Safety Cooperation Commission to monitor compliance and grant exceptions 2. A Joint AI Safety Laboratory to conduct R&D in AI safety to minimize existential risks and develop new AI models with no more than 10^{25} FLOPS	1. An Executive Council initially composed of the United States and the People's Republic of China 2. A Coalition Technology Body to coordinate activities of the parties, modify thresholds, and create an AI Technique Whitelist (with specified allowed AI methods and techniques)

^a Miotti, “An International Treaty to Implement a Global Compute Cap for Advanced Artificial Intelligence.”

^b TAISC.org, “Treaty on Artificial Intelligence Safety and Cooperation.”

^c Scher et al., “An International Agreement to Prevent the Premature Creation of Artificial Superintelligence.”

Notes

¹ Winter-Levy and Leicht, “The AI Divide: How U.S.-Chinese Competition Could Leave Most Countries Behind.”

² Stanford Institute for Human-Centered AI, *The 2025 AI Index Report*.

³ White House, *Winning the Race: America’s AI Action Plan*.

⁴ AlShebli et al., “China and the U.S. Produce More Impactful AI Research When Collaborating Together.”

⁵ Jamali and Chia, “Trump Gives Nvidia Green Light to Sell Advanced AI Chips to China.”

⁶ Kerrigan, Grek, and Mazarr, *The United States and China—Designing a Shared Future: The Potential for Track 2 Initiatives to Design an Agenda for Coexistence*.

⁷ Werbach, “U.S.-China AI Cooperation Under Trump 2.0.”

⁸ There is also one bilateral agreement on AI. In a November 2024 meeting, U.S. President Joe Biden and Chinese President Xi Jinping affirmed the need to address the risks of AI systems, improve AI safety and international cooperation, promote AI for the good of all, and maintain human control over the decision to use nuclear weapons. See White House, “Readout of President Joe Biden’s Meeting with President Xi Jinping of the People’s Republic of China,” for more details.

⁹ United Kingdom Department for Science, Innovation & Technology, “The Bletchley Declaration by Countries Attending the AI Safety Summit, 1–2 November 2023.”

¹⁰ United Nations General Assembly, “Seizing the Opportunities of Safe, Secure and Trustworthy Artificial Intelligence Systems for Sustainable Development.”

¹¹ United Nations General Assembly, “Enhancing International Cooperation on Capacity-Building of Artificial Intelligence.”

¹² White House, *Winning the Race: America’s AI Action Plan*.

¹³ Embassy of the People’s Republic of China in the Kingdom of Norway, “Global AI Governance Initiative.”

¹⁴ These proposed agreements represent the totality of proposed agreements identified by the author at the time of writing (spring 2026).

¹⁵ Miotti, “An International Treaty to Implement a Global Compute Cap for Advanced Artificial Intelligence.”

¹⁶ TAISC.org, “Treaty on Artificial Intelligence Safety and Cooperation.”

¹⁷ Scher et al., “An International Agreement to Prevent the Premature Creation of Artificial Superintelligence.”

¹⁸ See, for example, Uuk et al., “A Taxonomy of Systemic Risks from General-Purpose AI,” and Field, “Why Do Experts Disagree on Existential Risk and P(doom)? A Survey of AI Experts.”

¹⁹ The qualifier “somewhat” is necessary to highlight that although these treaties preempted disaster on Antarctica and in outer space, they were negotiated and agreed to in the aftermath of the destruction and nuclear employment of World War II.

²⁰ Antarctic Treaty; Treaty on Principles Governing the Activities of States in the Exploration and Use of Outer Space, Including the Moon and Other Celestial Bodies.

²¹ Atomic Heritage Foundation, “The Franck Report.”

²² U.S. Department of Defense, *Department of Defense International Science and Technology Engagement Strategy: A Unified Approach to Strengthen Alliances and Attract New Partners*.

²³ These efforts could leverage ongoing Trump administration initiatives, such as Pax Silica, and global mechanisms, such as the Summit on Responsible AI in the Military Domain.

²⁴ The Eisenhower administration supported the International Geophysical Year for its alignment with defense and economic priorities, achievement of international cooperation in the peaceful uses of atomic energy, and demonstration of U.S. leadership. For more details, see National Science Foundation, “The International Geophysical Year.”

²⁵ White House, “Readout of President Joe Biden’s Meeting with President Xi Jinping of the People’s Republic of China.”

²⁶ These two pathways are grounded in historical precedent. For example, President Dwight Eisenhower’s Atoms for Peace initiative sought to safely distribute the benefits of nuclear power. Simultaneously, nuclear powers sought to tightly restrain global proliferation of nuclear weapons.

²⁷ The activities in this appendix were co-written with perplexity.ai.

²⁸ One lesson from the coronavirus pandemic was that if the incident implicated a great power, then it could be immediately politicized and spawn new geopolitical tensions.

²⁹ The 1972 Basic Principles agreement between U.S. President Richard Nixon and Soviet General Secretary Leonid Brezhnev provided the political cover, conceptual framing, momentum, and normative language for later arms control agreements. Although it was legally nonbinding, the agreement eased détente-era negotiations and shaped rhetoric about crisis avoidance. For more details, see American Presidency Project, “Text of the ‘Basic Principles of Relations Between the United States of America and the Union of Soviet Socialist Republics.’”

References

AlShebli, Bedoor, Shahan Ali Memon, James A. Evans, and Talal Rahwan, “China and the U.S. Produce More Impactful AI Research When Collaborating Together,” *Scientific Reports*, Vol. 14, November 19, 2024.

American Presidency Project, “Text of the ‘Basic Principles of Relations Between the United States of America and the Union of Soviet Socialist Republics,’” webpage, undated. As of February 4, 2026: <https://www.presidency.ucsb.edu/documents/text-the-basic-principles-relations-between-the-united-states-america-and-the-union-soviet>

Antarctic Treaty, December 1, 1959. As of February 4, 2026: https://documents.ats.aq/keydocs/vol_1/vol1_2_AT_Antarctic_Treaty_e.pdf

Atomic Heritage Foundation, “The Franck Report,” webpage, undated. As of April 9, 2026: <https://ahf.nuclearmuseum.org/ahf/key-documents/franck-report/>

Embassy of the People’s Republic of China in the Kingdom of Norway, “Global AI Governance Initiative,” position paper, October 18, 2023. As of April 9, 2026: <https://no.china-embassy.gov.cn/eng/lcvt/lcwj/202401/P020240112008151194499.pdf>

Field, Severin, “Why Do Experts Disagree on Existential Risk and P(doom)? A Survey of AI Experts,” arXiv, arXiv:2502.14870v1, January 25, 2025.

Horowitz, Michael C., and Lauren A. Kahn, “Nuclear Non-Proliferation Is the Wrong Framework for AI Governance,” *AI Frontiers*, June 27, 2025. As of February 4, 2026: <https://ai-frontiers.org/articles/nuclear-non-proliferation-is-the-wrong-framework-for-ai-governance>

Jamali, Lily, and Osmond Chia, “Trump Gives Nvidia Green Light to Sell Advanced AI Chips to China,” BBC News, December 8, 2025.

Kerrigan, Amanda, Lydia Grek, and Michael J. Mazarr, *The United States and China—Designing a Shared Future: The Potential for Track 2 Initiatives to Design an Agenda for Coexistence*, RAND Corporation, RR-A2850-1, 2023. As of March 20, 2026: https://www.rand.org/pubs/research_reports/RRA2850-1.html

Miotti, Andrea, “An International Treaty to Implement a Global Compute Cap for Advanced Artificial Intelligence,” arXiv, arXiv:2311.10748v2, September 28, 2025.

National Science Foundation, “The International Geophysical Year,” June 2, 1954.

Scher, Aaron, David Abecassis, Peter Barnett, and Brian Abeyta, “An International Agreement to Prevent the Premature Creation of Artificial Superintelligence,” arXiv, arXiv:2511.10783, November 17, 2025. As of February 4, 2026: <https://arxiv.org/pdf/2511.10783>

Stanford Institute for Human-Centered AI, *The 2025 AI Index Report*, undated. As of April 8, 2026: <https://hai.stanford.edu/ai-index/2025-ai-index-report>

Strategic Arms Reduction Treaty (START I), July 31, 1991.

TAISC.org, “Treaty on Artificial Intelligence Safety and Cooperation,” webpage, undated. As of February 4, 2026: <https://taisc.org/taisc>

Treaty Banning Nuclear Weapon Tests in the Atmosphere, in Outer Space and Under Water (Partial Test Ban Treaty), August 5, 1963.

Treaty on Principles Governing the Activities of States in the Exploration and Use of Outer Space, Including the Moon and Other Celestial Bodies (Outer Space Treaty), December 19, 1966. As of February 4, 2026: <https://www.unoosa.org/oosa/en/ourwork/spacelaw/treaties/outerspacetreaty.html>

Treaty on the Non-Proliferation of Nuclear Weapons (Non-Proliferation Treaty), July 1, 1968.

United Kingdom Department for Science, Innovation & Technology, “The Bletchley Declaration by Countries Attending the AI Safety Summit, 1–2 November 2023,” webpage, updated February 13, 2025. As of February 4, 2026:

<https://www.gov.uk/government/publications/ai-safety-summit-2023-the-bletchley-declaration/the-bletchley-declaration-by-countries-attending-the-ai-safety-summit-1-2-november-2023>

United Nations General Assembly, “Seizing the Opportunities of Safe, Secure and Trustworthy Artificial Intelligence Systems for Sustainable Development,” Draft Resolution A/78/L.49, March 11, 2024. As of February 4, 2026:

<https://docs.un.org/en/A/78/L.49>

United Nations General Assembly, “Enhancing International Cooperation on Capacity-Building of Artificial Intelligence,” Resolution A/RES/78/311, July 5, 2024. As of February 4, 2026:

<https://docs.un.org/en/A/RES/78/311>

U.S. Department of Defense, *Department of Defense International Science and Technology Engagement Strategy: A Unified Approach to Strengthen Alliances and Attract New Partners*, October 2020.

Uuk, Risto, Carlos Ignacio Gutierrez, Daniel Guppy, Lode Lauwaert, Atoosa Kasirzadeh, Lucia Velasco, Peter Slattery, and Carina Prunkl, “A Taxonomy of Systemic Risks from General-Purpose AI,” arXiv, arXiv:2412.07780v1, November 24, 2024.

Werbach, Kevin, “U.S.-China AI Cooperation Under Trump 2.0,” Perry World House, November 24, 2025.

White House, “Readout of President Joe Biden’s Meeting with President Xi Jinping of the People’s Republic of China,” webpage, November 16, 2024. As of February 4, 2026:

<https://bidenwhitehouse.archives.gov/briefing-room/statements-releases/2024/11/16/readout-of-president-joe-bidens-meeting-with-president-xi-jinping-of-the-peoples-republic-of-china-3/>

White House, *Winning the Race: America’s AI Action Plan*, July 2025. As of April 8, 2026:

<https://www.whitehouse.gov/wp-content/uploads/2025/07/Americas-AI-Action-Plan.pdf>

Winter-Levy, Sam, and Anton Leicht, “The AI Divide: How U.S.-Chinese Competition Could Leave Most Countries Behind,” *Foreign Affairs*, February 10, 2026.

Acknowledgments

Many thanks to my RAND colleagues for their initial discussions, suggestions, and support, especially Joel Predd, Salman Ahmed, Matan Chorev, and Matt Chesson. I am deeply appreciative of Shanshan Mei and Daniel Remler for their thoughtful reviews of early drafts that helped strengthen this paper.

About the Author

Jared Mondschein is a physical scientist at RAND. His research focuses on the intersection of AI and U.S. foreign policy. From 2023 to 2025, Jared served on detail to the U.S. State Department's Office of the Special Envoy for Critical and Emerging Technology. He has a Ph.D. in chemistry.

About This Paper

This paper aims to inform artificial intelligence (AI) contingency planning under the premise that preserving strategic optionality for future U.S.-China engagement serves U.S. interests. While current geopolitical competition, deep mutual distrust, and the strategic stakes involved may make bilateral AI cooperation neither feasible nor sensible in the near term, this analysis examines what frameworks and mechanisms would be required if circumstances changed. Understanding how cooperation could work—even if it proves unnecessary or unattainable—can help policymakers preserve flexibility in an uncertain technological landscape. This paper is intended for policymakers in the United States and other countries and for academics in the AI research policy community.

Center for the Geopolitics of Artificial General Intelligence

RAND Global and Emerging Risks is a division of RAND that delivers rigorous and objective public policy research on the most consequential challenges to civilization and global security. This work was undertaken

by the division's Center for the Geopolitics of Artificial General Intelligence (AGI), which is committed to helping decisionmakers understand, anticipate, and prepare to navigate the national security and geopolitical implications of AGI. The center convenes leading technologists, strategists, economists, political scientists, and outside experts to consider the feasibility and effectiveness of prospective AGI-enabled capabilities; the domestic and international implications of their use; and the strategies and policies that governments, businesses, and civil society could adopt to respond to new realities. For more information, visit www.rand.org/geopolitics-of-agi.

Funding

This effort was independently initiated and conducted within the Center for the Geopolitics of Artificial General Intelligence using income from operations and gifts from RAND supporters, including philanthropic gifts made or recommended by DALHAP Investments Ltd., Ergo Impact, Founders Pledge, Charlottes och Fredriks Stiftelse, Good Ventures, Longview, and Coefficient Giving. RAND donors and grantors have no influence over research findings or recommendations.

RAND is a research organization that develops solutions to public policy challenges to help make communities throughout the world safer and more secure, healthier and more prosperous. RAND is nonprofit, nonpartisan, and committed to the public interest.

Research Integrity

Our mission to help improve policy and decisionmaking through research and analysis is enabled through our core values of quality and objectivity and our unwavering commitment to the highest level of integrity and ethical behavior. To help ensure our research and analysis are rigorous, objective, and nonpartisan, we subject our research publications to a robust and exacting quality-assurance process; avoid both the appearance and reality of financial and other conflicts of interest through staff training, project screening, and a policy of mandatory disclosure; and pursue transparency in our research engagements through our commitment to the open publication of our research findings and recommendations, disclosure of the source of funding of published research, and policies to ensure intellectual independence. For more information, visit www.rand.org/about/research-integrity.

RAND's publications do not necessarily reflect the opinions of its research clients and sponsors. **RAND**® is a registered trademark.

Limited Print and Electronic Distribution Rights

This publication and trademark(s) contained herein are protected by law. This representation of RAND intellectual property is provided for noncommercial use only. Unauthorized posting of this publication online is prohibited; linking directly to its webpage on rand.org is encouraged. Permission is required from RAND to reproduce, or reuse in another form, any of its research products for commercial purposes. For information on reprint and reuse permissions, please visit www.rand.org/about/publishing/permissions.

For more information on this publication, visit www.rand.org/t/PEA4825-1.

© 2026 RAND Corporation



www.rand.org